

# CS 461 / ECE 422

## Discussion #3

# AppSec: Checkpoint 2

Paul Murley – pmurley2@illinois.edu

9/11/2019

# Overview

- Buffer Overflow Basics
- Integer Overflows
- Exploits with Variable Stack Position
- Linked List Exploits
- Return-Oriented Programming (ROP)
- Format String Attacks
- Callback Shell

# Logistics

- MP1 Checkpoint 2 Due Wednesday, 9/19 @ 6:00pm
  - No late submissions accepted
- MP1 Checkpoint 1 is now graded
  - New branch called "AppSec\_cp1\_grades"
- MP1 Checkpoint 1 Regrade Requests
  - Due THIS MONDAY (9/16) @ 11:59pm
- MP2 (Web Security) Released Wednesday, 9/19
- Office Hours
  - Every weeknight from 5:00pm-7:00pm (starting 9/3)
  - Room: Siebel 4405

# Buffer Overflow Basics

- Altering control flow of a target program to cause it to do what attacker wants
- 1. Identify a code pointer that is eventually loaded into PC
- 2. Overwrite code pointer (memory write vulnerability)
- 3. Divert PC to code that will do useful work (for attacker)
- *In Aleph One attack:*
  - Code pointer: return address on stack
  - Memory write vulnerability: buffer overflow
  - Divert PC to shellcode in stack buffer

[https://www.eecs.umich.edu/courses/eecs588/static/stack\\_smashing.pdf](https://www.eecs.umich.edu/courses/eecs588/static/stack_smashing.pdf)

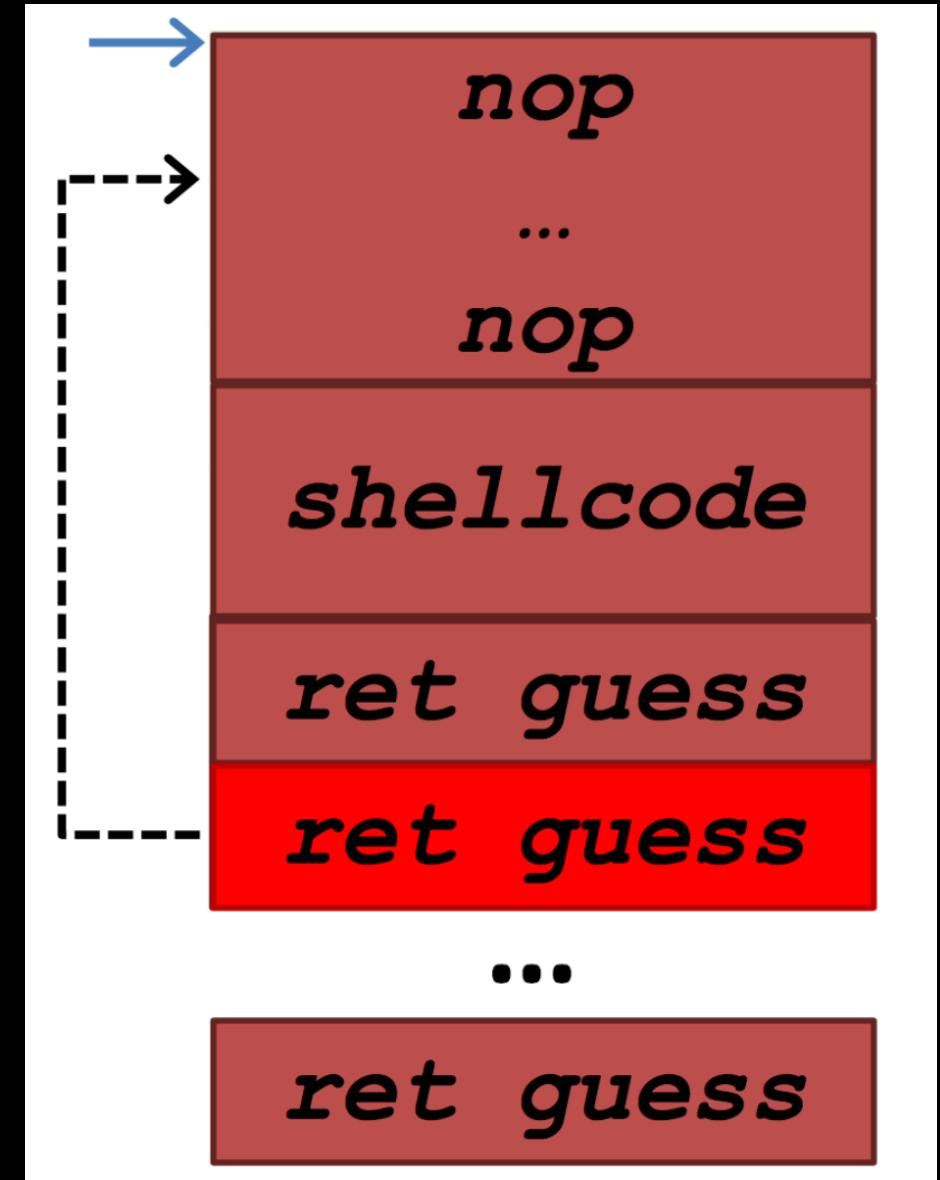
# Integer Overflows

- In 32-bit systems, there are maximum values for integers:
  - MAX\_INT is 2,147,483,647 (0x7FFFFFFF)
  - MAX\_UNSIGNED\_INT is 4,294,967,295 (0xFFFFFFFF)
- What happens when we overflow the maximum value, through addition or multiplication?

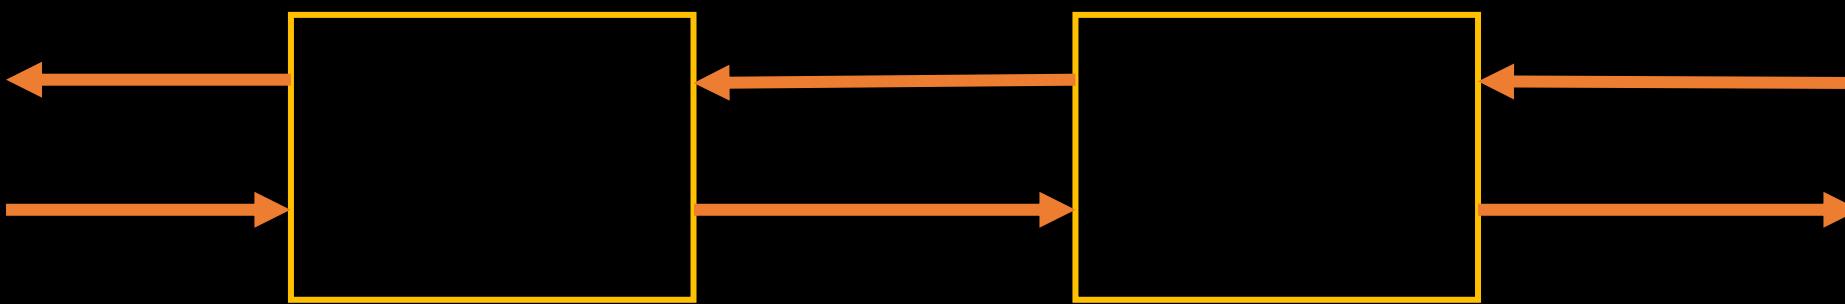
# Variable Stack Position

## ASLR (Address Space Layout Randomization):

A technique on modern operating systems that changes the memory layout of a program on every run, to make the development of exploits more difficult



# Linked List Exploitation



```
typedef struct node {
    struct node *next;
} Node;

add_lineage(Node* grandparent, Node* parent, Node* child) {
    grandparent->next = parent;
    grandparent->next->next = child;
}

main() {
    node* a = create_node()
    node* b = create_node();
    node* c = create_node();

    add_lineage(a,b,c);
}
```

```
typedef struct node {
    struct node *next;
} Node;

add_lineage(Node* grandparent, Node* parent, Node* child) {
    grandparent->next = parent;
    grandparent->next->next = child;
}

main() {
    node* a = create_node()
    node* b = create_node();
    node* c = create_node();

    add_lineage(a,b,c);
}
```

How can we write 0xdeadbeef into address 0xbffe1234?

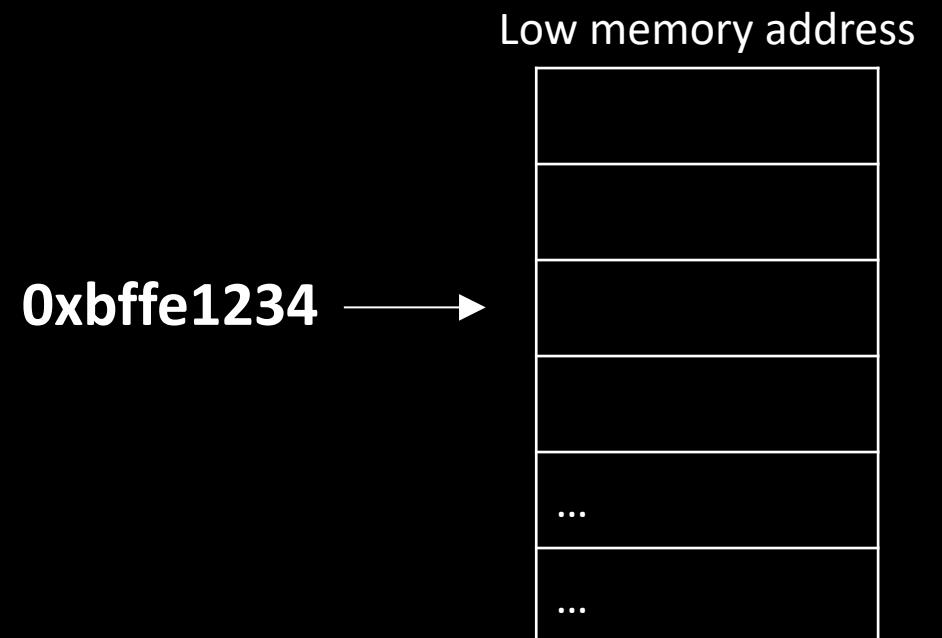
```
typedef struct node {  
    struct node *next;  
} Node;
```

```
add_lineage(Node* grandparent, Node* parent, Node* child) {  
    grandparent->next = parent;  
    grandparent->next->next = child;  
}
```

```
main () {  
    node* a = create_node ()  
    node* b = create_node () ;  
    node* c = create_node () ;  
  
    add_lineage(a,b,c) ;  
}
```

How can we write 0xdeadbeef into address 0xbffe1234?

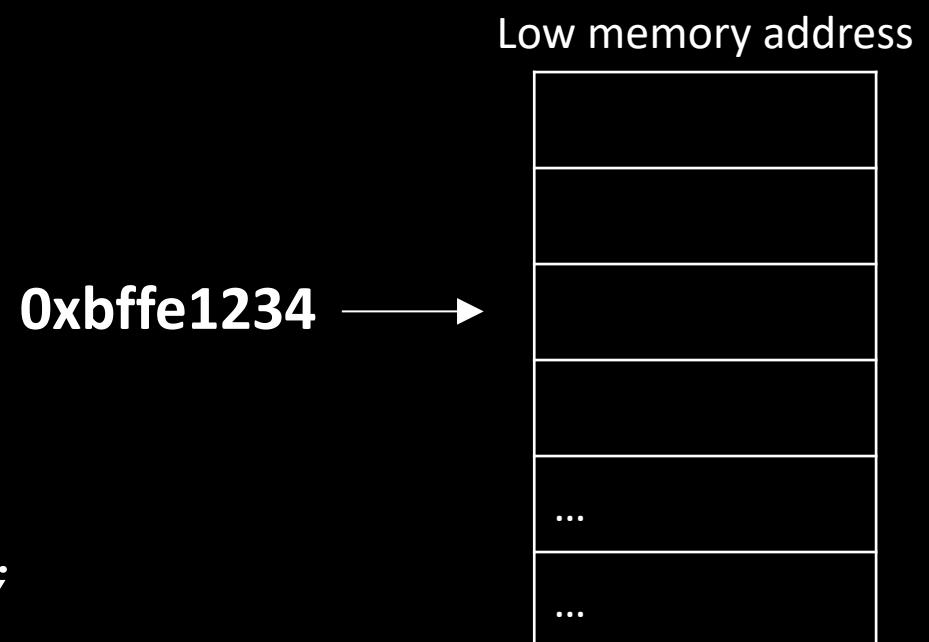
```
typedef struct node {  
    struct node *next;  
} Node;  
  
add_lineage(Node* grandparent, Node* parent, Node* child) {  
    grandparent->next = parent;  
    grandparent->next->next = child;  
}  
  
main () {  
    node* a = create_node ()  
    node* b = create_node () ;  
    node* c = create_node () ;  
  
    add_lineage(a,b,c) ;  
}
```



How can we write 0xdeadbeef into address 0xbffe1234?

```
typedef struct node {  
    struct node *next;  
} Node;
```

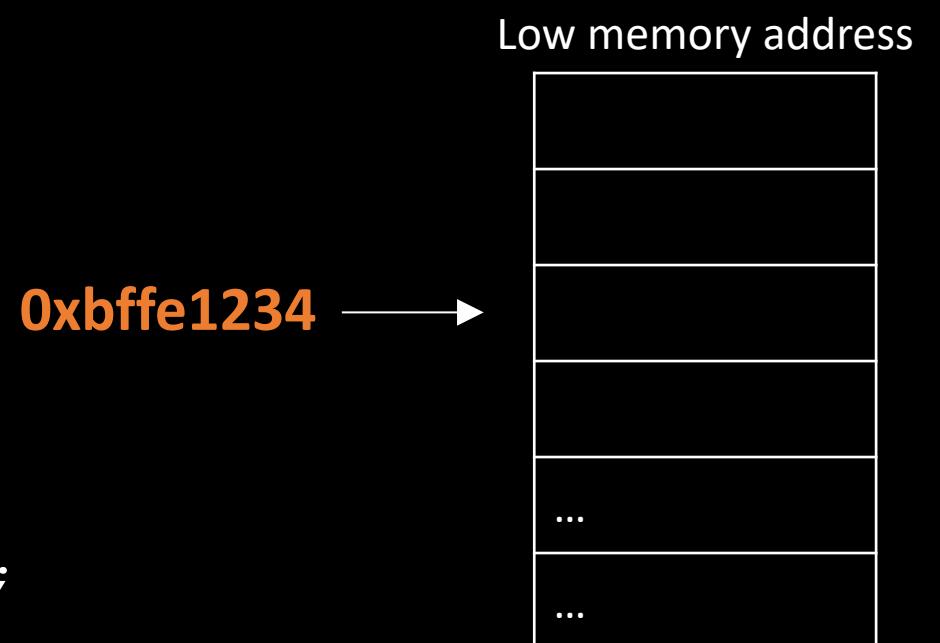
```
add_lineage(Node* grandparent, Node* parent, Node* child) {  
    grandparent->next = parent;  
    grandparent->next->next = child;  
}  
  
main () {  
    node* a = create_node ()  
    node* b = create_node () ;  
    node* c = create_node () ;  
  
    add_lineage(a, 0xbffe1234, 0xdeadbeef) ;  
}
```



How can we write 0xdeadbeef into address 0xbffe1234?

```
typedef struct node {  
    struct node *next;  
} Node;
```

```
add_lineage(Node* grandparent, Node* parent, Node* child) {  
    grandparent->next = parent;  
    grandparent->next->next = child;  
}  
  
main () {  
    node* a = create_node ()  
    node* b = create_node () ;  
    node* c = create_node () ;  
  
    add_lineage(a, 0xbffe1234, 0xdeadbeef) ;  
}
```

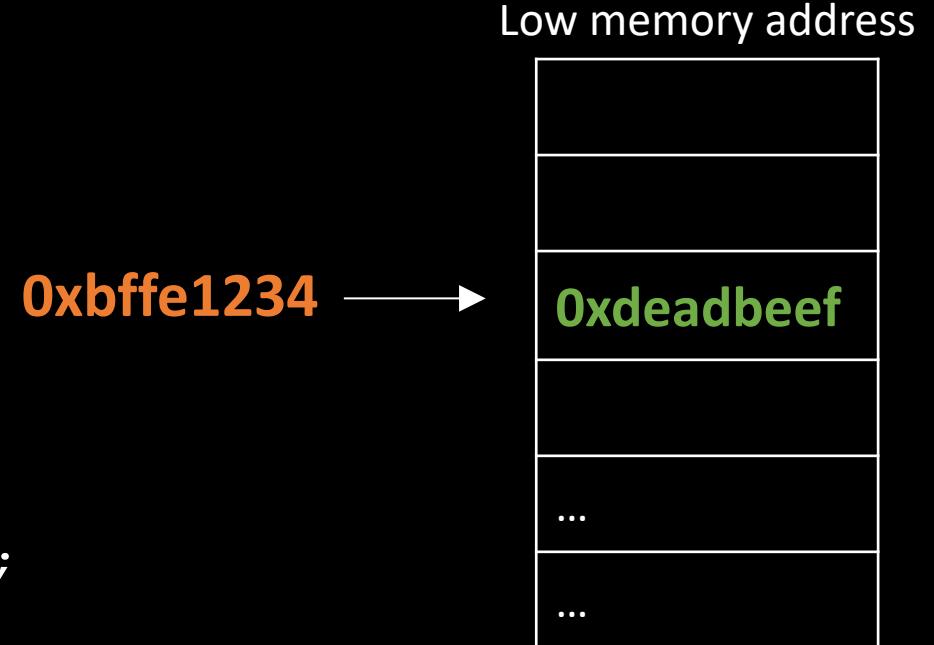


How can we write **0xdeadbeef** into address **0xbffe1234**?

```
typedef struct node {  
    struct node *next;  
} Node;
```

```
add_lineage(Node* grandparent, Node* parent, Node* child) {  
    grandparent->next = parent;  
    grandparent->next->next = child;  
}
```

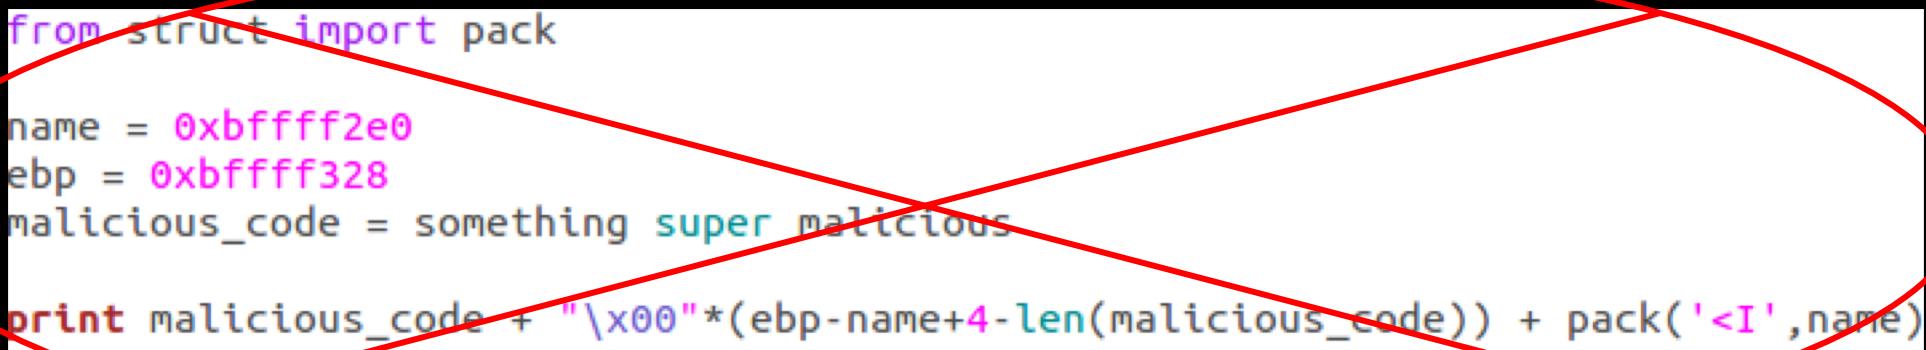
```
main () {  
    node* a = create_node ();  
    node* b = create_node ();  
    node* c = create_node ();  
  
    add_lineage(a, 0xbffe1234, 0xdeadbeef);  
}
```



How can we write **0xdeadbeef** into address **0xbffe1234**?

# Data Execution Prevention (DEP)

- Why do we need to be able to execute code on the stack?
  - Answer: Usually, there is no good reason why we should do this
- Defense Idea: Let's just disallow all data sections of memory from executing! This is called DEP or “Write XOR Execute”



```
from struct import pack

name = 0xfffff2e0
ebp = 0xfffff328
malicious_code = something super malicious

print malicious_code + "\x00"*(ebp-name+4-len(malicious_code)) + pack('<I',name)
```

# Return-Oriented Programming

- Idea: rather than running code that we inject, let's try to use code that already exists on the system
- General approach: Find “gadgets” in existing code. Use those gadgets in a sequence to execute assembly instructions at will
  - Gadgets:
    - Are usually 2-5 instructions long
    - End in ‘ret’
    - Accomplish some small task (move one register to another, push a value, pop a value into a register, zero out a register, etc.)

# Return-Oriented Programming

Disassembly of section .init:

```
080481c0 <_init>:  
 80481c0: 53                      push   %ebx  
 80481c1: 83 ec 08                sub    $0x8,%esp  
 80481c4: e8 00 00 00 00          call   80481c9 <_init+0x9>  
 80481c9: 5b                      pop    %ebx  
 80481ca: 81 c3 2b 6e 0a 00      add    $0xa6e2b,%ebx  
 80481d0: 8b 83 fc ff ff ff      mov    -0x4(%ebx),%eax  
 80481d6: 85 c0                  test   %eax,%eax  
 80481d8: 74 05                  je    80481df <_init+0x1f>  
 80481da: e8 21 7e fb f7          call   0 <__libc_tsd_LOCALE>  
 80481df: e8 bc 0c 00 00          call   8048ea0 <frame_dummy>  
 80481e4: e8 67 cd 07 00          call   80c4f50 <__do_global_ctors_aux>  
 80481e9: 83 c4 08                add    $0x8,%esp  
 80481ec: 5b                      pop    %ebx  
 80481ed: c3                      ret
```

# Return-Oriented Programming

Disassembly of section .init:

```
080481c0 <_init>:  
 80481c0: 53                      push   %ebx  
 80481c1: 83 ec 08                sub    $0x8,%esp  
 80481c4: e8 00 00 00 00          call   80481c9 <_init+0x9>  
 80481c9: 5b                      pop    %ebx  
 80481ca: 81 c3 2b 6e 0a 00      add    $0xa6e2b,%ebx  
 80481d0: 8b 83 fc ff ff ff      mov    -0x4(%ebx),%eax  
 80481d6: 85 c0                  test   %eax,%eax  
 80481d8: 74 05                  je     80481df <_init+0x1f>  
 80481da: e8 21 7e fb f7          call   0 <__libc_tsd_LOCALE>  
 80481df: e8 bc 0c 00 00          call   8048ea0 <frame_dummy>  
 80481e4: e8 67 cd 07 00          call   80c4f50 <__do_global_ctors_aux>  
 80481e9: 83 c4 08                add    $0x8,%esp  
 80481ec: 5b                      pop    %ebx  
 80481ed: c3                      ret
```

# Return-Oriented Programming

- Idea: rather than running code that we inject, let's try to use code that already exists on the system
- General approach: Find “gadgets” in existing code. Use those gadgets in a sequence to execute assembly instructions at will

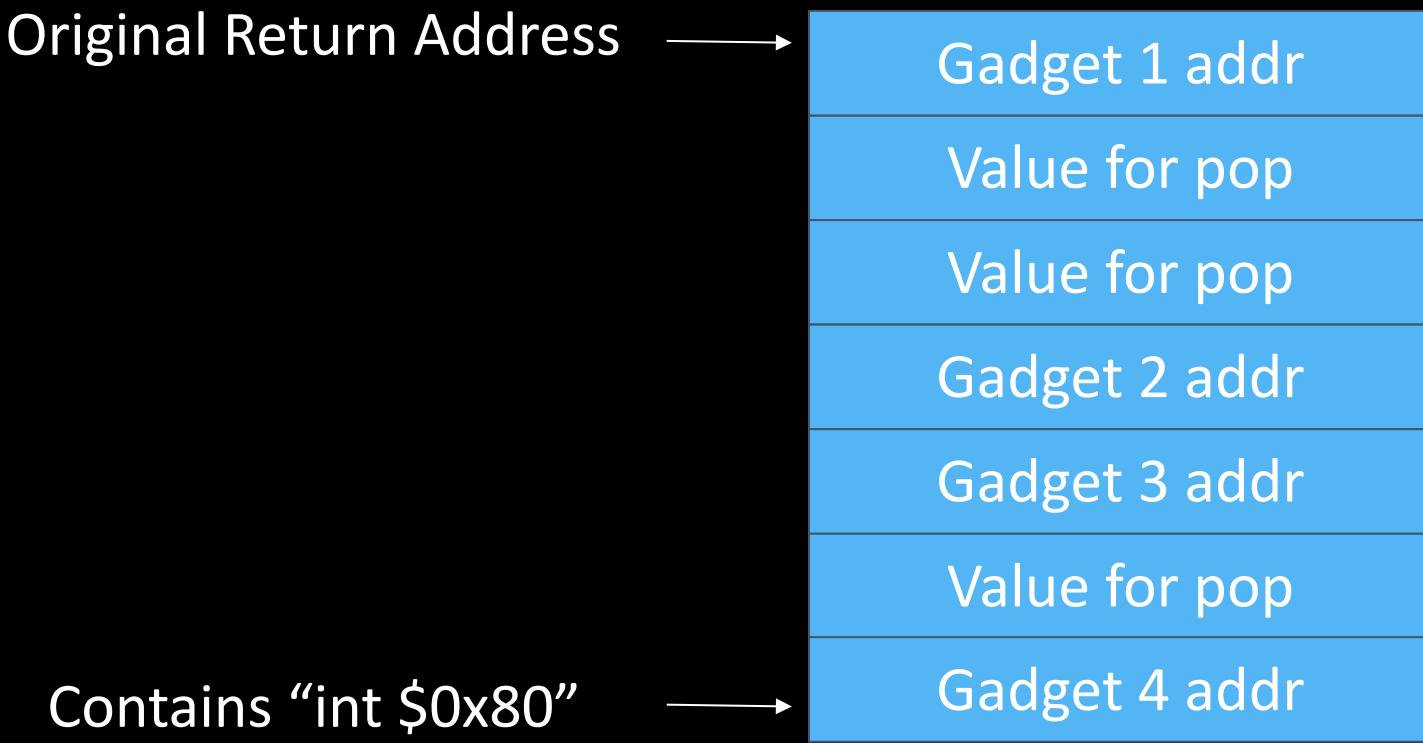
# Return-Oriented Programming

- Idea: rather than running code that we inject, let's try to use code that already exists on the system
- General approach: Find “gadgets” in existing code. Use those gadgets in a sequence to execute assembly instructions at will
- What are the possible challenges with this approach?

# Gadget Chaining

- For each gadget, keep track of your stack.
- If there's a “pop” instruction, prepare something on the stack to pop.
- If there's a “push” instruction, you don't need to prepare value for the next pop.
- Prepare the address of the beginning of next gadget (“ret” will pop that address into %eip, and two gadgets are chained together)
- The last gadget might end with “int \$0x80”. In that case, we don't need to prepare value on the stack for that

# Setting Up the Stack for ROP



# ROP Example

Goal: Put 0xdeafbeef into %eax

Overwritten Original Return Address



0x8057360

8057360:	5a	→	pop	%edx
8057361:	59		pop	%ecx
8057362:	5b		pop	%ebx
8057363:	c3		ret	

8055060:	8b 01	mov	(%ecx),%eax
8055062:	89 02	mov	%eax,(%edx)
8055064:	89 d0	mov	%edx,%eax
8055066:	c3	ret	

# ROP Example

Goal: Put 0xdeadbeef into %eax

Overwritten Original Return Address

8057360:	5a	→	pop	%edx
8057361:	59		pop	%ecx
8057362:	5b		pop	%ebx
8057363:	c3		ret	



0x8057360
0xdeadbeef (%edx)

8055060:	8b 01	mov	(%ecx),%eax
8055062:	89 02	mov	%eax,(%edx)
8055064:	89 d0	mov	%edx,%eax
8055066:	c3	ret	

# ROP Example

Goal: Put 0xdeadbeef into %eax

Overwritten Original Return Address

8057360:	5a	pop %edx
8057361:	59	pop %ecx
8057362:	5b	pop %ebx
8057363:	c3	ret

→ → → →

0x8057360
0xdeadbeef (%edx)
Garbage Value (%ecx)

8055060:	8b 01	mov (%ecx),%eax
8055062:	89 02	mov %eax,(%edx)
8055064:	89 d0	mov %edx,%eax
8055066:	c3	ret

# ROP Example

Goal: Put 0xdeadbeef into %eax

8057360:	5a	pop %edx
8057361:	59	pop %ecx
8057362:	5b	pop %ebx
8057363:	c3	ret



Overwritten Original Return Address

0x8057360
0xdeadbeef (%edx)
Garbage Value (%ecx)
Garbage Value (%ebx)

8055060:	8b 01	mov (%ecx),%eax
8055062:	89 02	mov %eax,(%edx)
8055064:	89 d0	mov %edx,%eax
8055066:	c3	ret

# ROP Example

Goal: Put 0xdeadbeef into %eax

8057360:	5a	pop %edx
8057361:	59	pop %ecx
8057362:	5b	pop %ebx
8057363:	c3	ret



Overwritten Original Return Address

0x8057360
0xdeadbeef (%edx)
Garbage Value (%ecx)
Garbage Value (%ebx)
0x8055064

8055060:	8b 01	mov (%ecx),%eax
8055062:	89 02	mov %eax,(%edx)
8055064:	89 d0	mov %edx,%eax
8055066:	c3	ret

# ROP Example

Goal: Put 0xdeadbeef into %eax

8057360:	5a	pop %edx
8057361:	59	pop %ecx
8057362:	5b	pop %ebx
8057363:	c3	ret



Overwritten Original Return Address

0x8057360
0xdeadbeef (%edx)
Garbage Value (%ecx)
Garbage Value (%ebx)
0x8055064
Next Gadget Address

8055060:	8b 01	mov (%ecx),%eax
8055062:	89 02	mov %eax,(%edx)
8055064:	89 d0	mov %edx,%eax
8055066:	c3	ret

## 1.2.9 (ROP) Tips

- Make sure your gadgets do no contain null characters, including in the addresses!
- Write down your goals first, then look for gadgets to help you achieve each goal
  - objdump -d ./program > program.txt
- For a deeper understanding, read the paper provided in the handout
- Be creative with gadgets (example: Put zero in %eax):

## 1.2.9 (ROP) Tips

- Make sure your gadgets do no contain null characters, including in the addresses!
- Write down your goals first, then look for gadgets to help you achieve each goal
  - objdump -d ./program > program.txt
- For a deeper understanding, read the paper provided in the handout
- Be creative with gadgets (example: Put zero in %eax):

0x80481ec

pop %eax  
inc %eax  
ret

Original Return Address →

0x80481ec

0xffffffff

Next Gadget

# Format String Attacks

```
void vulnerable(char *arg) {  
    char buf[2048];  
    strncpy(buf, arg, sizeof(buf));  
    printf(buf);  
}  
  
int _main(int argc, char **argv) {  
    if (argc != 2) {  
        fprintf(stderr, "Error: need a command-line argument\n");  
        return 1;  
    }  
    vulnerable(argv[1]);  
    return 0;  
}
```

# Format String Attacks

- Prototype answer: malicious\_code + padding + ADDR1 + ADDR2 + "%00000x%00\$hn%00000x%00\$hn"
- Your job: Find the right numbers
- Formats to understand:
  - %8x – prints 8 characters on the screen (doesn't matter what the chars are)
    - printf("%8x")      output: e39ed9a8
  - %n – writes to memory the number of characters printed
    - printf("AAAA%n", &i)    output: "AAAA", i = 4
  - 6\$ – '6\$' explicitly addresses the 6th parameter on the stack
    - printf ("%6\$d", 6, 5, 4, 3, 2, 1);    output: 1
  - %h – length modifier. Specifies that length is sizeof(short)
    - printf ("%hu", 0x10000);    output: 0
  - printf("%8x%6\$hn") –

# Format String Attacks

Goal: overwrite return address to point to buf that contains shellcode

Proto-answer: print malicious\_code + padding + ADDR1 + ADDR2 + "%00000x%04\$hn%00000x%05\$hn"

So, if your malicious code is at address 0xbfff1234:

- 1.print total of 0x1234 characters first, write to return address
- 2.print total of 0xbfff characters, write to return address+#  
(think about what the offset is)

# Callback Shell

- What is it?
- Why might an attacker want this functionality?

# Callback Shell

```
int sockfd;
struct sockaddr_in addr;

addr.sin_family = AF_INET;
addr.sin_addr.s_addr =
    inet_addr(SERV_HOST_ADDR);
addr.sin_port = htons(SERV_TCP_PORT);

sockfd = socket(AF_INET, SOCK_STREAM, 0);
connect(sockfd, (struct sockaddr *) &addr,
          sizeof(serv_addr));
do_stuff(stdin, sockfd);
```

# Callback Shell - Tasks

- Exploit itself it relatively straightforward (similar to 1.2.4)
- You need to write your own shellcode!
  - Create a client
  - Connect it to a server/listener (on given host/port)
  - Redirect stdin/stdout/stderr to the connection
  - Invoke a shell through the connection
- Recommendation: Write the code in x86 assembly and then translate it into hex once written

# Final Reminders

- MP1 Checkpoint 2 Due Wednesday, 9/19 @ 6:00pm
  - No late submissions accepted
- MP1 Checkpoint 1 Regrade Requests
  - Due THIS MONDAY (9/16) @ 11:59pm
- Office Hours
  - Every weeknight from 5:00pm-7:00pm (starting 9/3)
  - Room: Siebel 4405
- Contact
  - Paul Murley
  - [pmurley2@illinois.edu](mailto:pmurley2@illinois.edu)
  - CSL 445