

# Lecture 13: Web & Cloud Wrap-Up

Professor Adam Bates CS 461 / ECE 422 Fall 2019

## Goals for Today



- Learning Objectives:
  - Understand how cross-site scripting attacks violate the same origin policy (Web Wrap-Up)



- Explore why containers "do not contain"
- Announcements, etc:
  - Midterm October 9th, 7pm, 1404 Siebel
  - Grade distributions for MP1 checkpoints will be released after regrade requests are processed.
  - MP2 Checkpoint #1: Due Sept 25 at 6pm
  - MP2 Checkpoint #2: Due Oct 7 at 6pm



**Reminder**: Please put away devices at the start of class

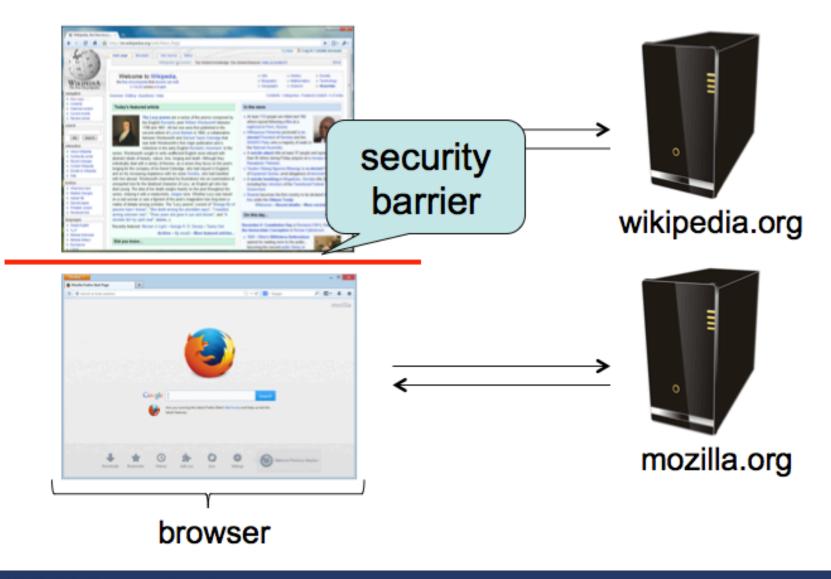
## Security on the web



- Risk #3: we don't want a malicious site to be able to spy on or tamper with my information or interactions with other websites
  - Browsing to evil.com should not let evil.com spy on my emails in Gmail or buy stuff with my Amazon account
- Defense: the same-origin policy
  - A security policy grafted on after-the-fact, and enforced by web browsers
  - Intuition: each web site is isolated from all others

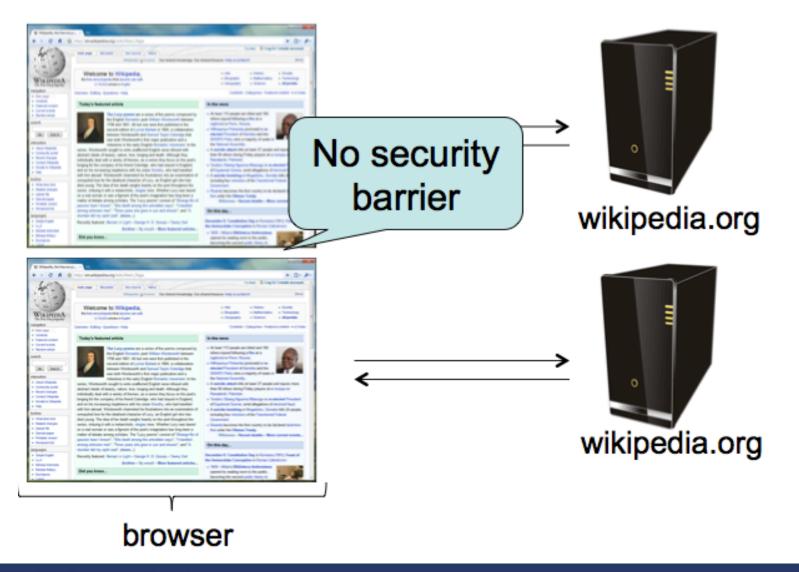


Each site is isolated from all others





Multiple pages from same site aren't isolated





- Granularity of protection: the origin
- Origin = protocol + hostname (+ port)



 Javascript on one page can read, change, and interact freely with all other pages from the same origin



- Browsers provide isolation for JS scripts via the Same Origin Policy (SOP)
- Simple version:
  - Browser associates web page elements (layout, cookies, events)
     with a given origin ≈ web server that provided the page/cookies in the first place
    - Identity of web server is in terms of its hostname, e.g., bank.com
- SOP = only scripts received from a web page's origin have access to page's elements
- XSS: Subverting the Same Origin Policy

#### Cross-site Request Forgery (CSRF)



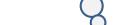
Suppose you log in to bank.com

POST /login?user=bob&pass=abc123 HTTP/1.1 Host: bank.com

HTTP/1.1 200 OK

Set-Cookie: login=fde874

• • • •



bank.com

fde874 = bob






#### Cross-site Request Forgery (CSRF)



GET /account HTTP/1.1

Host: bank.com

Cookie: login=fde874





HTTP/1.1 200 OK

••••

\$378.42



### Cross-site Request Forgery (CSRF)





Click me!!!

http://bank.com/transfer?to=badguy&amt=100

fde874 = bob

GET /transfer?to=badguy&amt=100 HTTP/1.1

Host: bank.com

Cookie: login=fde874

bank.com



HTTP/1.1 200 OK

• • • •

Transfer complete: -\$100.00



#### CSRF Defenses



- Need to "authenticate" each user action originates from our site
- One way: each "action" gets a token associated with it
  - On a new action (page), verify the token is present and correct
  - Attacker can't find token for another user,
     and thus can't make actions on the user's behalf

## CSRF Defenses



Pay \$25 to Joe:

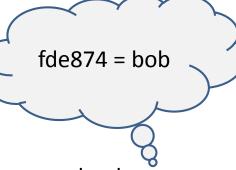
http://bank.com/transfer?to=joe&amt=25&token=8d64

<input type="hidden" name="token" value="8d64" />

HTTP/1.1 200 OK

Set-Cookie: token=8d64

• • • •



bank.com



GET /transfer?to=joe&amt=25&token=8d64 HTTP/1.1

Host: bank.com

Cookie: login=fde874

HTTP/1.1 200 OK

....

Transfer complete: -\$25.00



# Cross-Site Scripting (XSS)



```
<?php
echo "Hello, " . $_GET["user"] . "!";</pre>
```

GET /?user=Bob HTTP/1.1



HTTP/1.1 200 OK

•••

Hello, Bob!



# Cross-Site Scripting (XSS)



```
<?php
echo "Hello, " . $_GET["user"] . "!";</pre>
```

GET /?user=<u>Bob</u> HTTP/1.1



HTTP/1.1 200 OK

• • •

Hello, <u>Bob</u>!



# Cross-Site Scripting (XSS)



```
<?php
echo "Hello, " . $_GET["user"] . "!";</pre>
```

http://vuln.com/ says:

XSS



GET /?user=<script>alert('XSS')</script> HTTP/1.1

HTTP/1.1 200 OK

• • •

Hello, <script>alert('XSS')</script>!





#### Web Review | Same-Origin Policy (SOP)



```
GET / HTTP/1.1
Host: facebook.com
                HTTP/1.1 200 OK
                <script>
                $.get('http://gmail.com/msgs.json',
                  function (data) { alert(data); }
                </script>
GET /msgs.json HTTP/1.1
```

(evil!) facebook.com



Host: gmail.com

HTTP/1.1 200 OK



{ new\_msgs: 3 }





## Cross-Site Scripting (XSS) Attack



GET / HTTP/1.1

Host: facebook.com





२.get('http://gmail.com/ msgs.json', function (data) palert(data); })

HTTP/1.1 200 OK

<iframe src="http://gmail.com/?user=<script> \$.get('http://gmail.com/msgs.json', form (data) { alert(data); }) </script>"></iframe>



GET /?user=<script>\$.get(' ... </script> HTTP/1.1

Host: gmail.com

HTTP/1.1 200 OK

Hello, <script>\$.get('http://gmail.com/ msgs.json',

unction (data) { alert(data); }) </script>

gmail.com



## Cross-Site Scripting (XSS) Attack



```
(evil!)
                                                                   facebook.com
                     Host: facebook.com
http://gmail.com/ says:
{ new_msgs: 3 }
                                GET /msgs.json HTTP/1.1
                     Host: gmail.com
                                                                   gmail.com
                         HTTP/1.1 200 OK
                          { new_msgs: 3 }
```

### XSS Defenses



- Make sure data gets shown as data, not executed as code!
- Escape special characters
  - Which ones? Depends what context your \$data is presented
    - Inside an HTML document? <div>\$data</div>
    - Inside a tag? <a href="http://site.com/\$data">
    - Inside Javascript code? var x = "\$data";
  - Make sure to escape every last instance!
- Frameworks can let you declare what's user-controlled data and automatically escape it



and now, cloud wrap-up

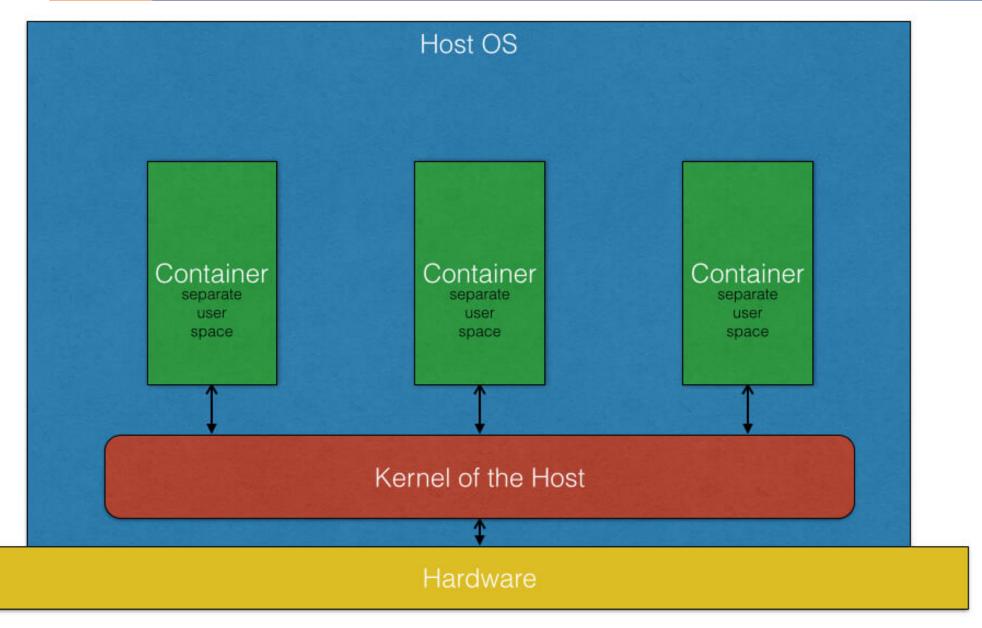
## Enter Containers



- Rather than virtualize both user space and kernel space... why not just 'virtualize' user space?
- Meets the needs of most customers, who don't require significant customization of the OS.
- Sometimes called 'operating system virtualization,' which is highly misleading...
- Running natively on host, containers enjoy bare metal performance without reliance on advanced virtualization support from hardware.

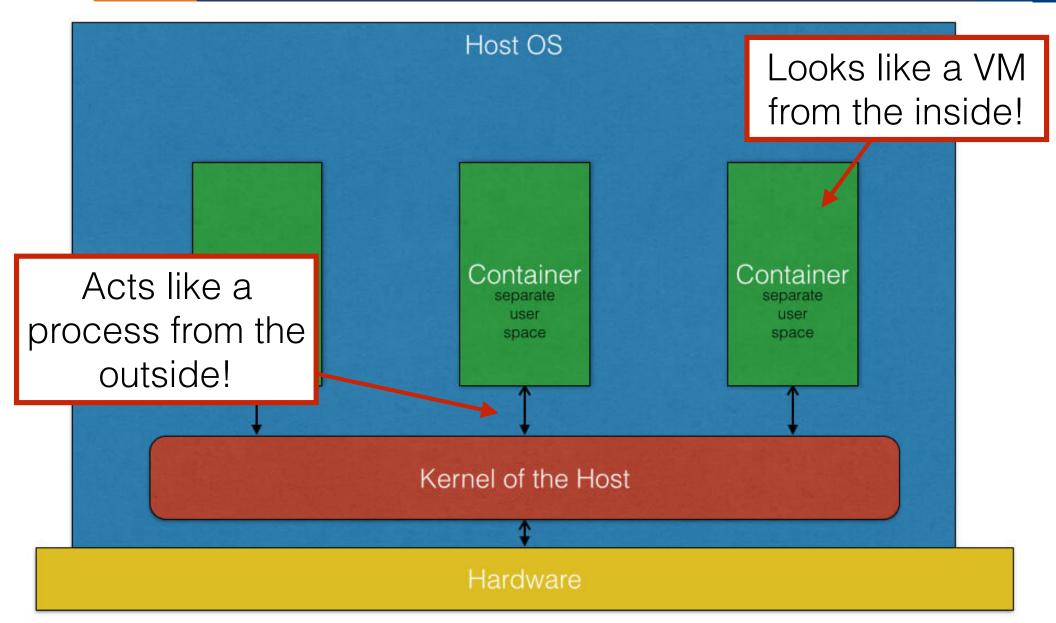
## Enter Containers





## Enter Containers





## OS Support for Containers



- Linux Containers (LXC):
  - chroot
  - Kernel Namespaces
    - PID, Network, User, IPC, uts, mount
  - cgroups for HW isolation
  - Security profiles and policies
    - Apparmor, SELinux, Seccomp

#### containers = chroot on steroids



- chroot changes the apparent root directory for a given process and all of its children
- An old idea! POSIX call dating back to 1979
- Not intended to defend against privileged attackers...
  they still have root access and can do all sorts of
  things to break out (like chroot'ing again)
- Hiding the true root FS isolates a lot; in \*nix, file abstraction used extensively.
- Does not completely hide processes, network, etc., though!

## Namespaces



- The key feature enabling containerization!
- Partition practically all OS functionalities so that different process domains see different things
- Mount (mnt): Controls mount points
- <u>Process ID (pid)</u>: Exposes a new set of process IDs distinct from other namespaces (i.e., the hosts)
- Network (net): Dedicated network stack per container; each interface present in exactly I namespace at a time.

•

## Namespaces



- The key feature enabling containerization!
- Partition practically all OS functionalities so that different process domains see different things
- Interprocess Comm. (IPC): Isolate processes from various methods of POSIX IPC.
  - e.g., no shared memory between containers!
- <u>UTS</u>: Allows the host to present different host/domain names to different containers.
- There's also a <u>User ID</u> (user) and <u>cgroup</u> namespace

## User Namespace



- Like others, can provide a unique UID space to the container.
- More nuanced though we can map UID 0 inside the container to UID 1000 outside; allows processes inside of container to think they're root.
- Enables containers to perform administration actions, e.g., adding more users, while remaining confined to their namespace.

#### cgroups



- Limit, track, and isolate utilization of hardware resources including CPU, memory, and disk.
- Important for ensuring QoS between customers! Protects against bad neighbors
- Operate at the namespace granularity, not per-process
- Features:
  - Resource limitation
  - Prioritization
  - Accounting (for billing customers!)
  - Control, e.g., freezing groups
- The cgroup namespace prevents containers from viewing or modifying their own group assignment

## Container Security?



"Containers do not contain." - Dan Walsh (SELinux contributor)

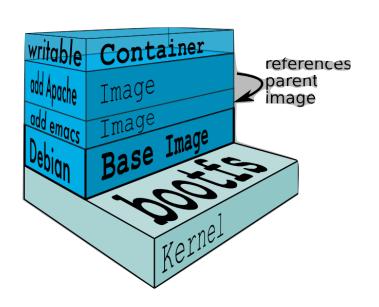
- In a nutshell, it's <u>real hard</u> to prove that every feature of the operating system is namespaced.
  - /sys? /proc? /dev? LKMs? kernel keyrings?
- Root access to any of these enables pwning the host
- Solution? Just don't forget about access control; SELinux now offers pretty good support for namespace labeling.
- SELinux and Namespaces actually synergize nicely; <u>much</u> easier to express a correct isolation policy over a coarse-grained namespace than, say, individual processes

#### How Docker fits in





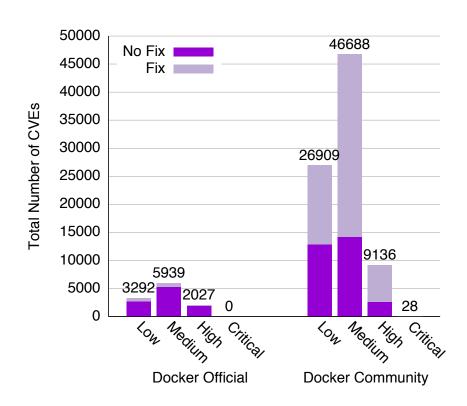
- Utilities that allow you to leverage (e.g.) LXC to build a portable, self-sufficient application using containers.
- Assures that all libraries and dependencies are packaged inside of a container image

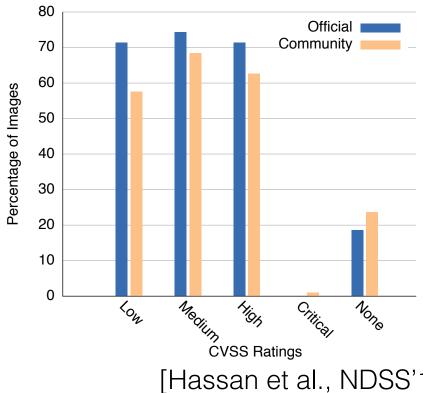


## Docker Attack Surface?



- Docker provides something analogous to an app market for building containers.
- Are the container images 'secure'?





[Hassan et al., NDSS'18]

#### Above the clouds...



- Container (~PaaS clouds) are strictly easier to manage than traditional laaS VMs.
- The era of Container hype has somewhat come and gone... containers still expose more flexibility than most users need!!
- Th hype now is about Function-as-a-Service cloud; (a.k.a. serverless computing)
- Individual programs/functions executed by invocation, great for event-driven stuff.
- Enabled by containers



# Serverless Computing



- Web Application is cached inside of FaaS cloud; no persistent architecture (virtual or otherwise)
- Web App components (Functions) are spawned and destroyed on a per-request basis.
- Function invocation is triggered by different event sources, e.g., web requests
- Pay by the millisecond for use
- Security considerations??

## Serverless Security



- Are functions secure?
- Standard attacks probably won't work, and if they do the function will shut down immediately
- Harder for malware to establish persistence, e.g., file system is read-only, container dies\*
  - \* it's supposed to, anyway...
- Harder to exfiltrate due to virtual network rules
  - reverse shell probably not available because exploitable function does not have Internet access

### FaaS Threat Vectors



- Probe with standard attacks (e.g., code injection) in the hopes of breaking something,
   identify location in workflow that will leak the presence of a crash (e.g., confirmation email)
- Can't drop to shell; how to establish persistence?

[Jones, CCC Congress'13]

## Container Reuse



- Functions execute inside a container
- For performance 'warm' containers are sometimes reused across multiple functions/customers
  - i.e., numerous explicit data channels (unlike laaS)
- Containers/Functions are cached in memory and relaunched if a request arrives in a short enough period of time
- Attacker can leverage this to simulate persistence

#### FaaS Threat Vectors



- Probe with standard attacks (e.g., code injection) in the hopes of breaking something,
   identify location in workflow that will leak the presence of a crash (e.g., confirmation email)
- Can't drop to shell; how to establish persistence?
  - Write malware or tool sets to /tmp
  - Issue periodic requests to keep container cached
- On AWS, customer access tokens are stored as environment variables in instance...
   potential for tokens to be wildly overprivileged if misconfigured.

```
{
    "AWS LAMBDA FUNCTION VERSION": "$LATEST",
    "LAMBDA_TASK_ROOT": "/var/task",
    "PATH": /usr/local/bin:/usr/bin/:/bin",
    "LD LIBRARY_PATH": "/lib64:/var/lib64:/var/runtime:/var/runtime/lib:/var/task:/var/task/lib",
    "LANG": "en_US.UTF-8",
    "AWS_LAMBDA_FUNCTION_NAME": "your-function-name",
    "AWS_REGION": "us-east-1",
    "AWS_REGION": "us-east-1",
    "AWS_SECON_TOOKEN": "FXXDYIZdEK3///////SFLKJBSKKLDJFLKJDFLSKJDFLSKJDFLSKJDFLKAJDSLKJHSGF",
    "AWS_SECON_TOOKEN": "FXXDYIZdEK3///////WELKSDJFLKABFDJa88asdf8asdfF==",
    "LAMBDA_RUNTIME_DIR": "/var/runtime",
    "AWS_LAMBDA_FUNCTION_MEMORY_SIZE": "512",
    "PYTHOMPATH": "/var/runtime",
    "AWS_LAMBDA_LOG_GROUP_NAME": "/aws/lambda/your-function-name",
    "AWS_LAMBDA_LOG_STREAM_NAME": "2016/12/13/[$LATEST]448bf3a99b754fe781006b5f6358b67b",
    "AWS_LAMBDA_LOG_STREAM_NAME": "2016/12/13/[$LATEST]448bf3a99b754fe781006b5f6358b67b",
    "AWS_LAMBDA_LOG_STREAM_NAME": "2016/12/13/[$LATEST]448bf3a99b754fe781006b5f6358b67b",
    "AWS_LAMBDA_LOG_STREAM_NAME": "2016/12/13/[$LATEST]448bf3a99b754fe781006b5f6358b67b",
    "AWS_LAMBDA_LOG_STREAM_NAME": "WMS_CAMBDA_LOG_STREAM_NAME": "WMMS_CAMBDA_LOG_STREAM_NAME": "WMS_CAMBDA_LOG_STREA
```

Use token-based authorizations to access data, then exfiltrate using cloud services

[Jones, CCC Congress'16]