



Lecture 32: Isolation & Authorization

Professor Adam Bates
CS 46I / ECE 422
Fall 2019

Goals for Today




- Learning Objectives:

- Define Isolation, its properties, and why is isolation important
- Understand at what levels isolation can be implemented
- Provide examples at each level



- Announcements, etc:

- Networking CP2: Due Nov 11
- **Wednesday is cancelled!**
 - ... not really, but no lecture or discussion
 - Next lecture is on Friday



Reminder: Please put away devices at the start of class

Security Properties



- Confidentiality?
 - Only trusted parties can read data
- Integrity?
 - Only trusted parties have modified data
- Authenticity?
 - Data originates from the correct party
- Availability?
 - Data is available to trusted parties when needed



Security Functions



- Define security functions over *principals* (e.g., users, programs, sysadmins)
- ... and also *entities* (e.g., files, network sockets, ipc)
- Authentication
 - How do we determine the identity of the principal?
- Authorization
 - Which principals are permitted to take what actions on which objects?
- Auditing
 - Record of (un) authorized actions that took place on the system for post-hoc diagnostics

Running untrusted code




- We often need to run buggy/untrusted code:
 - programs from untrusted Internet sites:
 - apps, extensions, plug-ins, codecs for media player
 - exposed applications: pdf viewers, outlook
 - legacy daemons: sendmail, bind
 - honeypots
- Goal: if application “misbehaves” \Rightarrow kill it

Approach: Confinement




- Confinement: ensure app cannot affect rest of system, or vice versa
- Can be implemented at many levels:
 - Hardware: run application on isolated hw (air gap)



Approach: Confinement



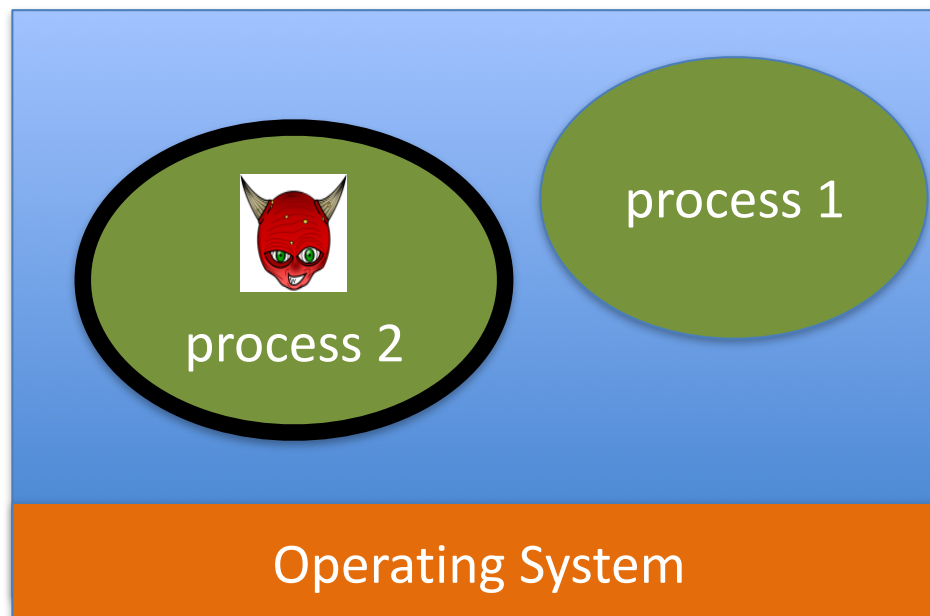
- Confinement: ensure app cannot affect rest of system, or vice versa
- Can be implemented at many levels:
 - Virtual machines: isolate OS's on a single machine



Approach: Confinement



- Confinement: ensure app cannot affect rest of system, or vice versa
- Can be implemented at many levels:
 - Process: System Call Interposition, Isolate a process in a single operating system



A old example: chroot



- Often used for “guest” accounts on ftp sites
- To use do: (must be root)

```
chroot /tmp/guest  
su guest
```

root dir “/” is now “/tmp/guest”
EUID set to “guest”

- Now “/tmp/guest” is added to file system accesses for applications in jail
- `open(“/etc/passwd”, “r”) ⇒`
`open(“/tmp/guest/etc/passwd” ,“r”)`
- application cannot access files outside of jail



- Problem: all utility progs (ls, ps, vi) must live inside jail
- jailkit project: auto builds files, libs, and dirs needed in jail env
 - jk_init: creates jail environment
 - jk_check: checks jail env for security problems
 - checks for any modified programs,
 - checks for world writable directories, etc.
 - jk_lsh: restricted shell to be used inside jail
- note: simple chroot jail does not limit network access


Escaping from jails



- Early escapes: relative paths
`open("../../etc/passwd", "r")` \Rightarrow
`open("/tmp/guest/../../etc/passwd", "r")`
- chroot should only be executable by root.
 - otherwise jailed app can do:
 - create dummy file `"/aaa/etc/passwd"`
 - run `chroot "/aaa"`
 - run `su root` to become root

(bug in Ultrix 4.0)

Virtual Machines



Example: **NSA NetTop**

single HW platform used for both classified and unclassified data

Why so popular now?



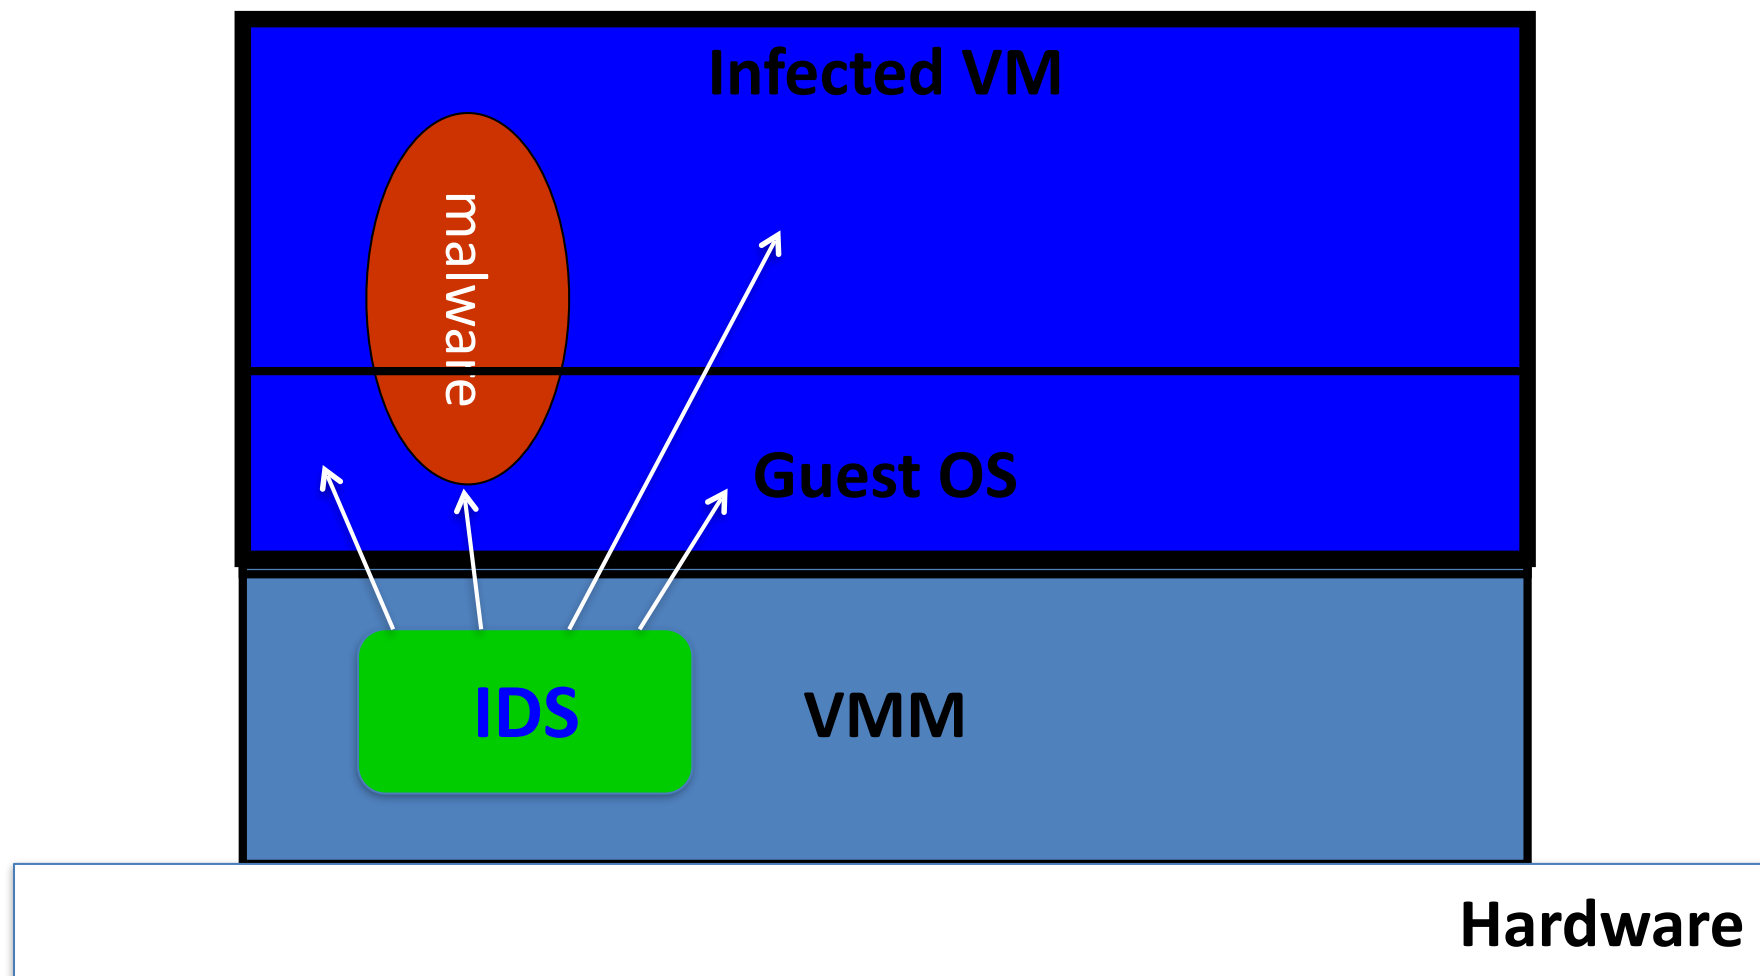
- VMs in the 1960's:
 - Few computers, lots of users
 - VMs allow many users to share a single computer
- VMs 1970's – 2000: non-existent
- VMs since 2000:
 - Too many computers, too few users
 - Print server, Mail server, Web server, File server, Database , ...
 - Wasteful to run each service on different hardware
 - More generally: VMs heavily used in cloud computing



- VMM Security assumption:
 - Malware can infect guest OS and guest apps
 - But malware cannot escape from the infected VM
 - Cannot infect host OS
 - Cannot infect other VMs on the same hardware
- Requires that VMM protect itself and is not buggy
 - VMM is much simpler than full OS
 - ... but device drivers run in Host OS



- Runs as part of OS kernel and user space process
 - Kernel root kit can shutdown protection system
 - Common practice for modern malware
- Standard solution: run IDS system in the network
 - Problem: insufficient visibility into user's machine
- Better: run IDS as part of VMM (protected from malware)
 - VMM can monitor virtual hardware for anomalies
 - VMI: Virtual Machine Introspection
 - Allows VMM to check Guest OS internals



Sample checks



- Stealth root-kit malware:
 - Creates processes that are invisible to “ps”
 - Opens sockets that are invisible to “netstat”
- 1. Lie detector check
 - Goal: detect stealth malware that hides processes and network activity
 - Method:
 - VMM lists processes running in GuestOS
 - VMM requests GuestOS to list processes (e.g. ps)
 - If mismatch: kill VM

Sample checks




- 2. Application code integrity detector
 - VMM computes hash of user app code running in VM
 - Compare to whitelist of hashes
 - Kills VM if unknown program appears
- 3. Ensure GuestOS kernel integrity
 - example: detect changes to `sys_call_table`
- 4. Virus signature detector
 - Run virus signature detector on GuestOS memory

Subvirt [King et al. 2006]



- Problem: Attackers can leverage VM isolation too
 - Once on victim machine, install a malicious VMM
 - Virus hides in VMM
 - Invisible to virus detector running inside VM





- ***Isolation is extremely coarse-grained:***
 - All or nothing access
 - Inappropriate for apps like a web browser
 - Needs read access to files outside jail (e.g. for sending attachments in Gmail)
- chroot does not prevent malicious apps from:
 - Accessing network and messing with other machines
 - Trying to crash host OS
- In practice, processes need to be able to cooperate



- Access control matrix
 - For every protected resource, list of who is permitted to do what
 - Example: for each file/directory, a list of permissions
 - Owner, group, world: read, write, execute
 - Setuid: program run with permission of principal who installed it
 - Smartphone: list of permissions granted each app

Crypto v. Access Control




- Two competing primitives for achieving security
- Both can be used to achieve assurances of data confidentiality, integrity, authenticity
- Crypto: “Encrypt the data for security!”
 - Key is authorization mechanism to assure confidentiality, integrity, authenticity properties.
- Access Control: “Label the data for security!”
 - Software includes authorization mechanism to assure confidentiality, integrity, authenticity properties according to policy.

Crypto v. Access Control



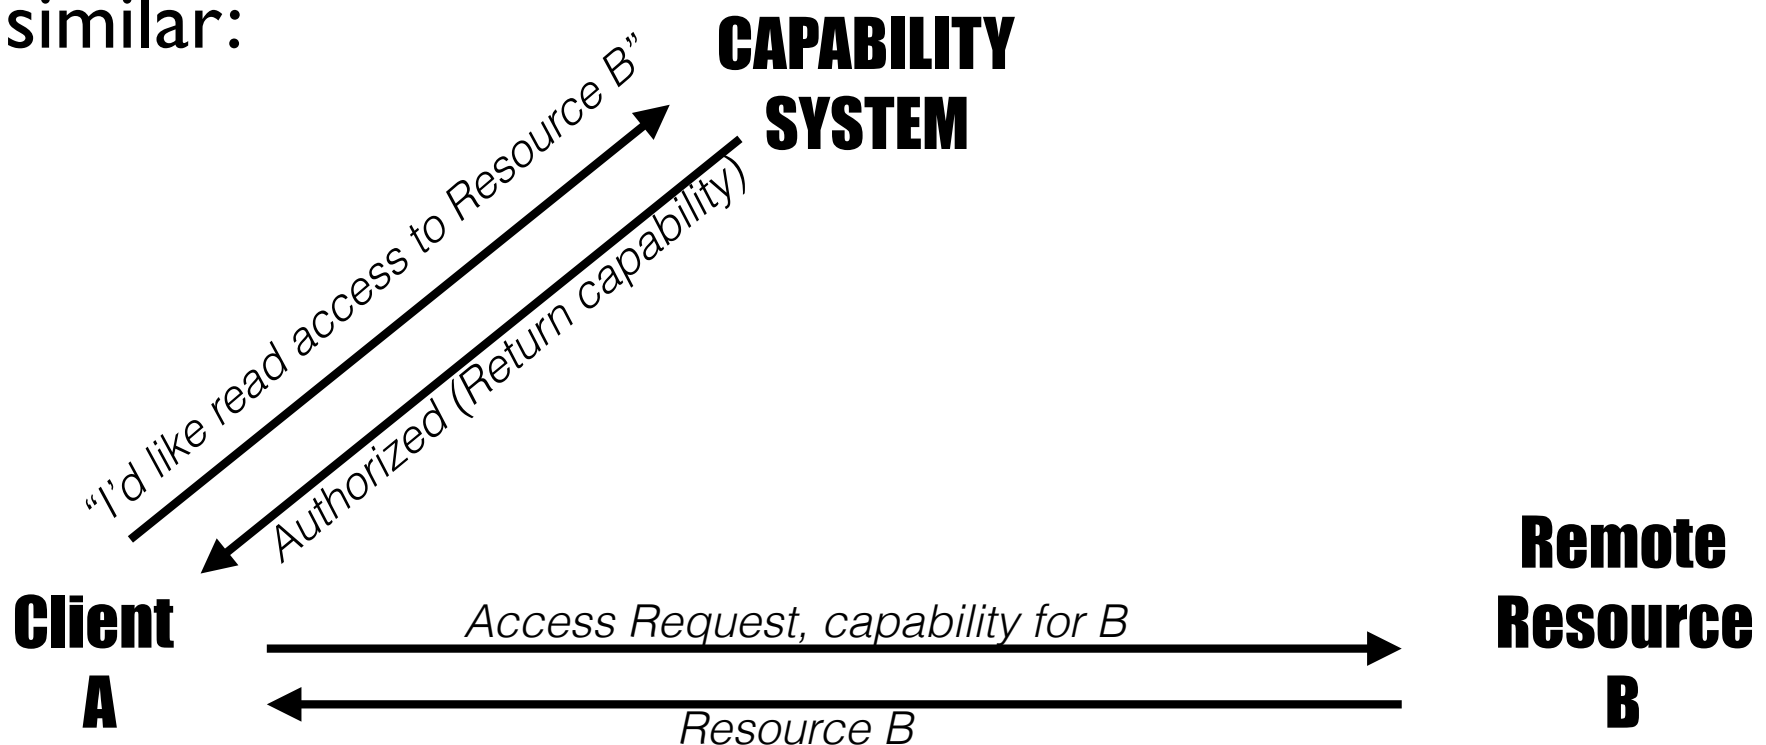
- Two competing primitives for achieving security
- Both can be used to achieve assurances of data confidentiality, integrity, authenticity
- Crypto- and AC- based architectures are often very similar:



Crypto v. Access Control




- Two competing primitives for achieving security
- Both can be used to achieve assurances of data confidentiality, integrity, authenticity
- Crypto- and AC- based architectures are often very similar:



Crypto v. Access Control



- Two competing primitives for achieving security
- Both can be used to achieve assurances of data confidentiality, integrity, authenticity
- Crypto- and AC- based architectures are often very similar:





- Which primitive should we use in operating systems?

Cryptography

- Easily supports multiple administrative domains
- Easily supports delegation (i.e., share key)
- Computationally expensive
- Requires key management and distribution mechanisms
- Requires key security

Access Control

- Struggles to support multiple administrative domains
- Easily supports centralized global security policy
- Computationally cheap
- Requires policy management mechanisms
- Requires reference monitor

Crypto v. Access Control



- Which primitive should we use in operating systems?

Cryptography

- Easily supports multiple administrative domains
- Easily supports delegation (i.e., share key)
- Computationally expensive
- Requires key management and distribution mechanisms
- Requires key security

Access Control

- Struggles to support multiple administrative domains
- Easily supports centralized global security policy
- Computationally cheap
- Requires policy management mechanisms
- Requires reference monitor

Linux has had broad crypto support since version 2.5, primarily for verifying signatures...

...but in the kernel, access control is king!

Access Control



- Determine whether a principal can perform a requested operation on a target object
- **Principal/Subject:** user, process, etc.
- **Operation/Action:** read, write, etc.
- **Object:** file, tuple, etc.

Protection Domains




- A computer system is a set of processes and objects
- A process operates within a protection domain
- A protection domain specifies the resources a process may access and the types of operations that may be invoked on the objects.
- The Principle of Least Privilege: The protection domain of a process should be as small as possible given the need of that process to accomplish its assigned task.

Protection Domains



Example of 3 protection domains:



What are domains based on in Linux?

Discretionary Access Control (DAC)



Find the access control!


```
batesa@sp19-cs423-adm:/dev$ ls -alh /dev
total 0
drwxr-xr-x 19 root root    4.0K Mar  6 06:48 .
drwxr-xr-x 23 root root    321 Mar  6 06:48 ..
crw----- 1 root root    10, 175 Feb 10 14:31 agpgart
crw----- 1 root root    10, 235 Feb 10 14:31 autofs
drwxr-xr-x 2 root root    380 Feb 10 14:31 block
drwxr-xr-x 2 root root     80 Feb 10 14:31 bsg
crw-rw---- 1 root disk   10, 234 Feb 10 14:31 btrfs-control
lrwxrwxrwx 1 root root      3 Feb 10 14:31 cdrom -> sr0
lrwxrwxrwx 1 root root      3 Feb 10 14:31 cdrw -> sr0
drwxr-xr-x 2 root root    3.3K Mar  6 06:48 char
crw----- 1 root root      5,  1 Feb 10 14:31 console
lrwxrwxrwx 1 root root     11 Feb 10 14:31 core -> /proc/kcore
drwxr-xr-x 6 root root    120 Mar  6 06:48 cpu
crw----- 1 root root    10,  59 Feb 10 14:31 cpu_dma_latency
crw----- 1 root root    10, 203 Feb 10 14:31 cuse
drwxr-xr-x 6 root root    120 Feb 10 14:31 disk
brw-rw---- 1 root disk  253,  0 Feb 10 14:31 dm-0
brw-rw---- 1 root disk  253,  1 Feb 10 14:31 dm-1
brw-rw---- 1 root disk  253,  2 Feb 10 14:31 dm-2
brw-rw---- 1 root disk  253,  3 Feb 10 14:31 dm-3
drwxr-xr-x 2 root root     80 Feb 10 14:31 dri
lrwxrwxrwx 1 root root      3 Feb 10 14:31 dvd -> sr0
crw----- 1 root root    10,  61 Feb 10 14:31 ecryptfs
crw-rw---- 1 root video   29,  0 Feb 10 14:31 fb0
lrwxrwxrwx 1 root root     13 Feb 10 14:31 fd -> /proc/self/fd
brw-rw---- 1 root disk      2,  0 Feb 10 14:31 fd0
```

- Owner or creator of resources specify which subjects have which access to those resources
- Commonly implemented in commercial products
- Access is managed by individual users, not a central security policy

Discretionary Access Control (DAC)



Access Mask defines permissions for User, Group, and Other



```
ba@p19-cs423-adm:/dev$ ls -alh /dev
total 0
drwxr-xr-x 19 root root 4.0K Mar  6 06:48 .
drwxr-xr-x 23 root root 321 Mar  6 06:48 ..
crw-rw-rw- 1 root root 10, 175 Feb 10 14:31 agpgart
crw-rw-rw- 1 root root 10, 235 Feb 10 14:31 autofs
drwxr-xr-x 2 root root 380 Feb 10 14:31 block
drwxr-xr-x 2 root root 80 Feb 10 14:31 bsg
crw-rw-rw- 1 root disk 10, 234 Feb 10 14:31 btrfs-control
lrwxrwxrwx 1 root root 3 Feb 10 14:31 cdrom -> sr0
lrwxrwxrwx 1 root root 3 Feb 10 14:31 cdrw -> sr0
drwxr-xr-x 2 root root 3.3K Mar  6 06:48 char
crw-rw-rw- 1 root root 5, 1 Feb 10 14:31 console
lrwxrwxrwx 1 root root 11 Feb 10 14:31 core -> /proc/kcore
drwxr-xr-x 6 root root 120 Mar  6 06:48 cpu
crw-rw-rw- 1 root root 10, 59 Feb 10 14:31 cpu_dma_latency
drwxr-xr-x 6 root root 10, 203 Feb 10 14:31 cuse
crw-rw-rw- 1 root root 120 Feb 10 14:31 disk
brw-rw-rw- 1 root disk 253, 0 Feb 10 14:31 dm-0
brw-rw-rw- 1 root disk 253, 1 Feb 10 14:31 dm-1
brw-rw-rw- 1 root disk 253, 2 Feb 10 14:31 dm-2
brw-rw-rw- 1 root disk 253, 3 Feb 10 14:31 dm-3
drwxr-xr-x 2 root root 80 Feb 10 14:31 dri
lrwxrwxrwx 1 root root 3 Feb 10 14:31 dvd -> sr0
crw-rw-rw- 1 root root 10, 61 Feb 10 14:31 ecryptfs
crw-rw-rw- 1 root video 29, 0 Feb 10 14:31 fb0
lrwxrwxrwx 1 root root 13 Feb 10 14:31 fd -> /proc/self/fd
brw-rw-rw- 1 root disk 2, 0 Feb 10 14:31 fd0
```

```
chmod u=rwx,g=rx,o=r myfile
chmod 754 myfile
```

<- Same thing

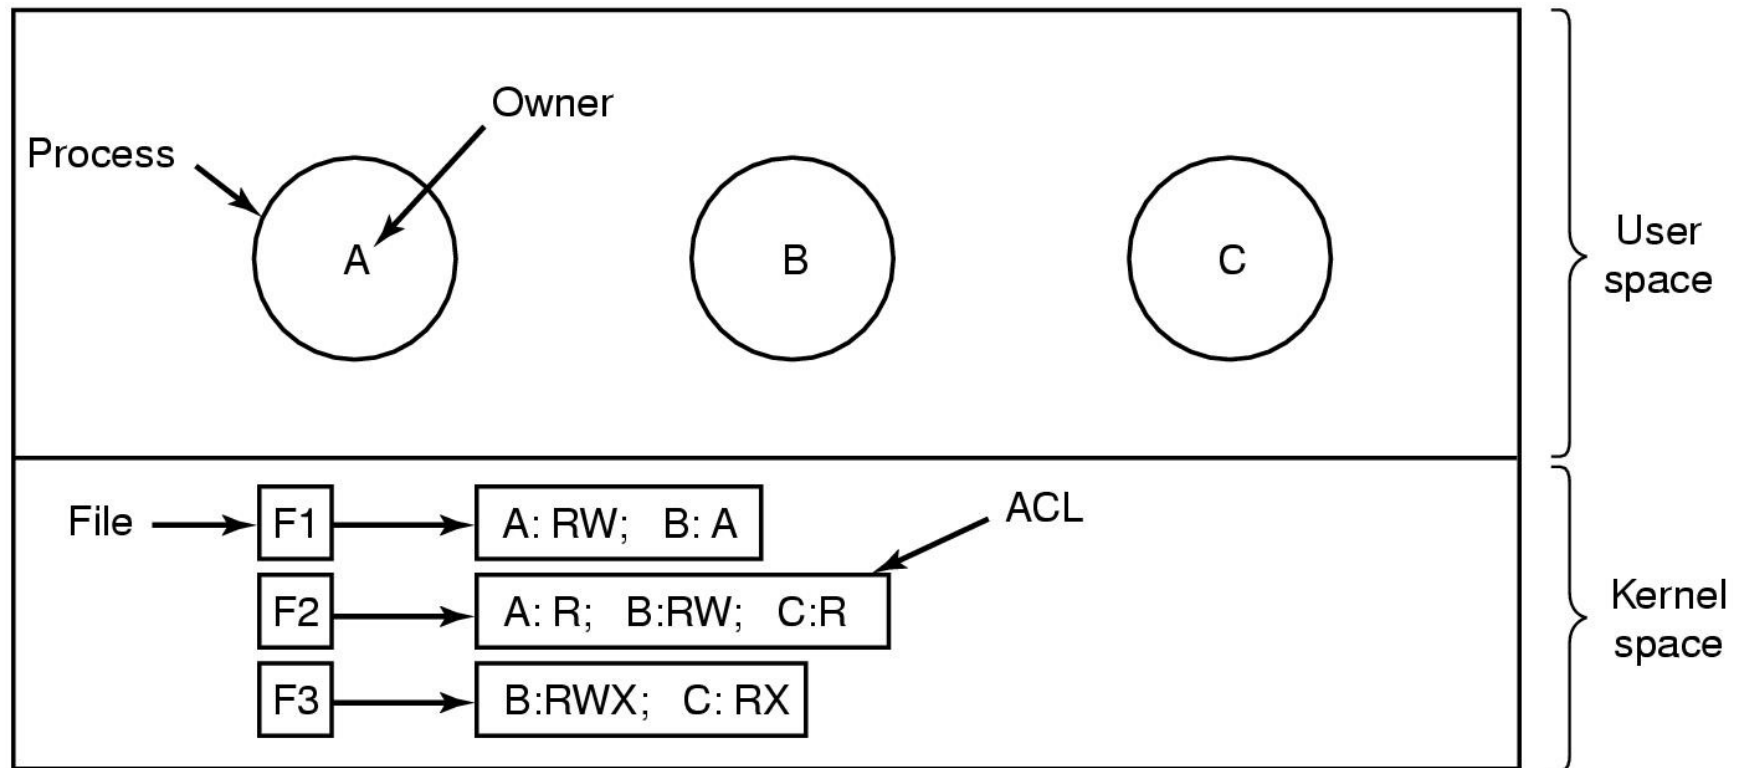
4 stands for "read",
2 stands for "write",
1 stands for "execute", and
0 stands for "no permission."

Access Control Lists (ACLs)



- Each column in access matrix specifies access for one Object.
- On invocation of a method R on an object O by a process running in a domain D ...
 - ... the access control list is searched to check whether D is allowed to perform method R on object O (e.g., allowed to read the file or execute the program)
- A default (e.g., “rest of the world”) can be associated with an access list so that any Domain not specified in the list can access the Objects using default methods.
- It is easy for the owner of the Object to grant access to another Domain or revoke access.
- ACL entries can be for individual users or for a group of users.

UNIX ACLs



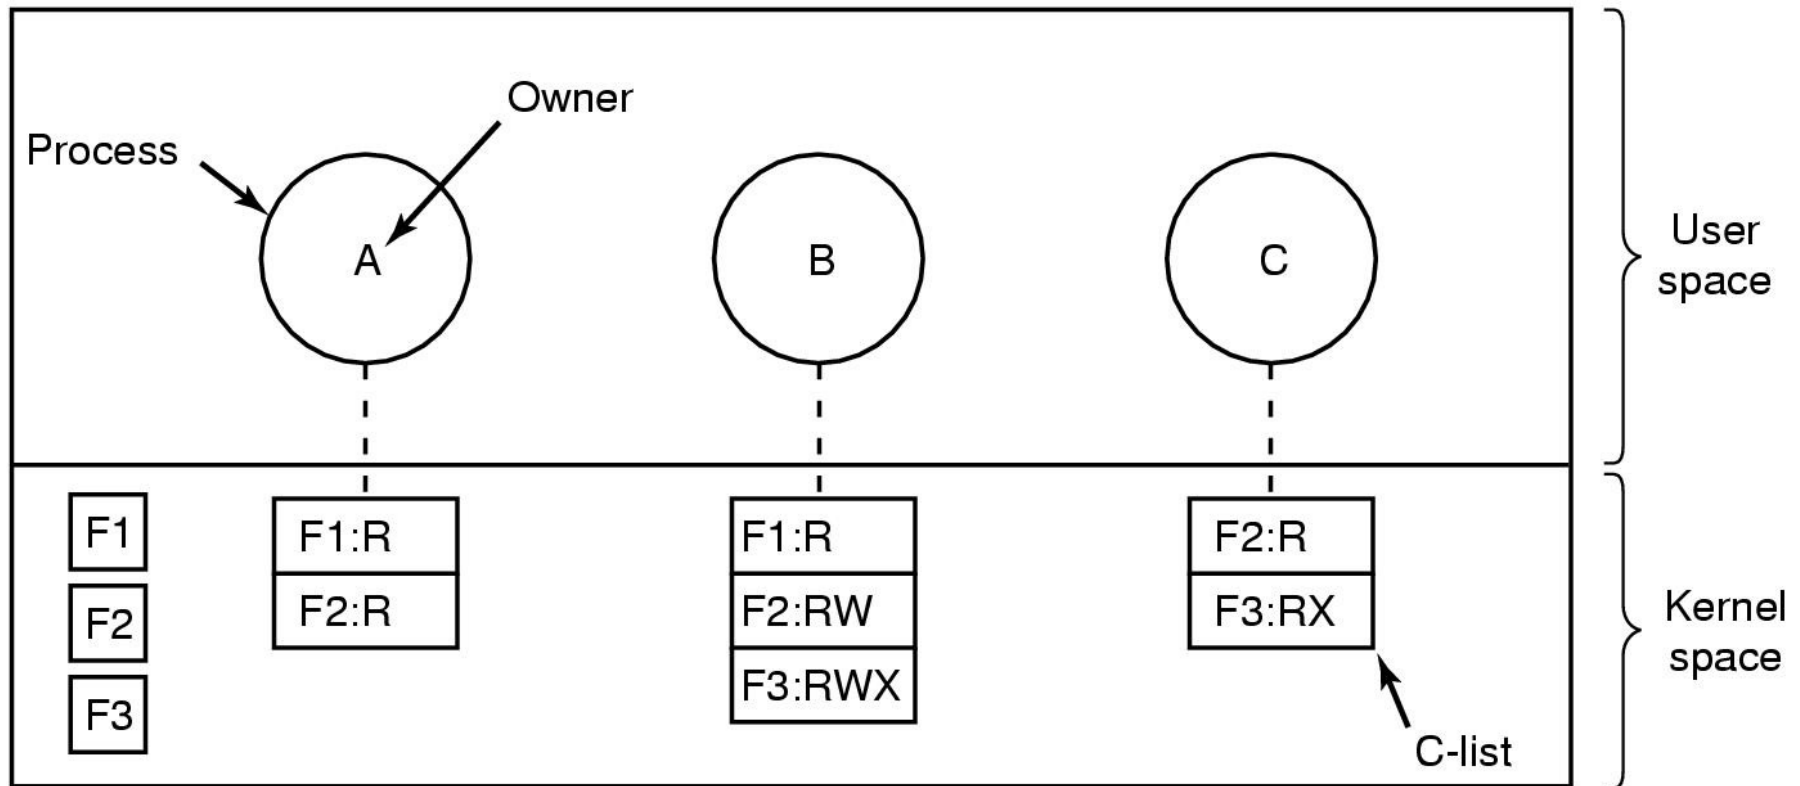
Use of access control lists to manage file access in UNIX

Capabilities



- An alternative access control mechanism
- Capability is an unforgeable ticket
 - Managed by OS
 - Can be passed from one process to another
- Permissive
- OS Mechanism (Reference monitor) checks ticket
 - Does not need to know the identity of the user/proc
- Can be used to partition superuser privilege
- Implementation: POSIX.1e capabilities

Capabilities



When capabilities are used, each process has a capability list.

Problems?



- What might go wrong with DAC or Capabilities?
 - Security is left to the discretion of subjects
 - Impossible to guarantee security of system
 - Security of system changes over time.
- Solution?
 - Mandatory Access Control: Operating system constrains the ability of subjects (even owners) to perform operations on objects according to a system-wide *security policy*.

To Learn More ...



- Books
 - Stallings and Brown, Chapter 12
 - Pfleeger and Pfleeger, Chapter 5
 - Goodrich and Tamassia, Chapter 3
 - Bishop, Chapter 16, 26, 29
- Papers
 - App Isolation: Get the Security of Multiple Browsers with Just One - Chen
 - Native Client: A Sandbox for Portable, untrusted x86 Native Code - Yee*
 - Innovative Instructions and Software Model for Isolated Execution - McKeen
 - The Security Architecture of the Chromium Browser – Barth*